

Recurrence Relations & The Master Theorem

Suhas Arehalli

COMP221 - Spring 2026

1 Analyzing Runtimes of Recursive Functions

Our techniques thus far taught us to analyze runtime by first computing growth functions (roughly) and then applying the appropriate simplifications to express that function in Big-O notation. While this technique is universal, we quickly get into trouble when analyzing certain kinds of recursive functions. Consider MERGESORT in Alg. 1.

Algorithm 1 Mergesort, with Merge's structure implied. See the MergeSort notes for details on Merge!

```
function MERGESORT(Array A[1...n])
    if N ≤ 1 then
        return A
    end if
    m ← ⌈ n / 2 ⌉
    L ← MERGESORT(A[1...m])
    R ← MERGESORT(A[m + 1...n])
    return MERGE(L, R)
end function
```

If we apply our standard tools, we see that the cost of a call to MERGESORT is a bunch of $\Theta(1)$ work, a $\Theta(n)$ call to MERGE, and... two recursive calls to MERGESORT, each on an input roughly half the size of the full array.

It might initially be confusing what to do with these recursive calls! If we don't know the runtime complexity of MERGESORT, we can't know what the recursive calls cost in time complexity and thus... can't determine the runtime complexity of MERGESORT. Our first new tool — recurrence relations — are meant just to get us proper notation to talk about this!

Let $T(n)$ denote the worst-case growth function for MERGESORT. We can express the runtime of MERGESORT as

$$T(n) = \begin{cases} 2T(\frac{n}{2}) + \Theta(n), & n > 1 \\ \Theta(1), & n \leq 1 \end{cases}$$

Note a few things about my notation here. First, it's a *recurrence relation*: we allow for $T(n)$ to appear on both sides (i.e., $T(n)$ is defined in terms of itself). This matches the recursive nature

of the algorithm! Additionally, note that (unlike Erickson) I set up the recurrence as a piecewise function, making explicit that there is both a recursive and base case. This is what I consider good analytic hygiene — this is tedious, but both technically correct and good at reminding you of the structure of the problem (we're trying to approach a base case!). Third, worth reminding you of our shorthand for Big-O: It is common to write, for instance, $\Theta(g(n))$ in the place of some function $f(n) \in \Theta(g(n))$. Finally, our cue to do more work, is to note that this is not a nice *closed-form* — it doesn't look like something we can, for instance, turn into Big-O notation! We need to write a non-recursive definition of $T(n)$!

1.1 Unrolling Recurrences

The first approach is to *unroll* the recursive part of the recurrence until you deduce a pattern. To unroll, we just substitute the equation into itself over and over, noting any patterns with respect to the number of substitutions.

$$\begin{aligned} T(n) &= 2T\left(\frac{n}{2}\right) + \Theta(n) \\ &= 2(2T\left(\frac{n}{4}\right) + \Theta(\frac{n}{2})) + \Theta(n) \\ &= 4T\left(\frac{n}{4}\right) + 2\Theta\left(\frac{n}{2}\right) + \Theta(n) \\ &= 4(2T\left(\frac{n}{8}\right) + \Theta(\frac{n}{4})) + 2\Theta(\frac{n}{2}) + \Theta(n) \\ &= 8T\left(\frac{n}{8}\right) + 4\Theta\left(\frac{n}{4}\right) + 2\Theta\left(\frac{n}{2}\right) + \Theta(n) \\ &= 2^3T\left(\frac{n}{2^3}\right) + 2^2\Theta\left(\frac{n}{2^2}\right) + 2^1\Theta\left(\frac{n}{2^1}\right) + 2^0\Theta\left(\frac{n}{2^0}\right) \end{aligned}$$

So after k substitutions we find

$$T(n) = 2^{k+1}T\left(\frac{n}{2^{k+1}}\right) + \sum_{i=0}^k 2^i\Theta\left(\frac{n}{2^i}\right)$$

And have a hope of getting rid of the T term on the right hand side — we simply need to choose k such that $T(\frac{n}{2^{k+1}})$ is a base case! This requires $\frac{n}{2^{k+1}} \leq 1$, Which solves out to $k \geq \log n - 1$, and so we want $k = \log n - 1$ and we can substitute $\Theta(1)$ for our call to T , getting

$$\begin{aligned} T(n) &= 2^{\log n}\Theta(1) + \sum_{i=0}^{\log n - 1} 2^i\Theta\left(\frac{n}{2^i}\right) \\ &= n\Theta(1) + \sum_{i=0}^{\log n - 1} 2^i\Theta\left(\frac{n}{2^i}\right) \\ &= \Theta(n) + \sum_{i=0}^{\log n - 1} \Theta(n) \\ &= \log n \cdot \Theta(n) = \Theta(n \log n) \end{aligned}$$

And we can conclude that MERGESORT is $\Theta(n \log n)$ as desired!
 We can do the same for something like Binary Search! Consider Alg. 2

Algorithm 2 Binary Search

```

function BINARYSEARCH(Array  $A[1 \dots n]$ , target  $e$ )
    if  $n == 0$  then
        return NULL
    end if
     $m \leftarrow \lceil \frac{n}{2} \rceil$ 
    if  $A[m] == e$  then
        return  $m$ 
    else if  $e \leq A[m]$  then
        return BINARYSEARCH( $A[1 \dots m - 1]$ ,  $e$ )
    else
        return BINARYSEARCH( $A[m + 1 \dots n]$ ,  $e$ )
    end if
end function
    
```

Whose recurrence relation for worst-case runtime will be

$$T(n) = \begin{cases} T\left(\frac{n}{2}\right) + \Theta(1), & N > 0 \\ \Theta(1), & N = 0 \end{cases}$$

And then we can unroll, getting

$$\begin{aligned} T(n) &= T\left(\frac{n}{2}\right) + \Theta(1) \\ &= \left(T\left(\frac{n}{4}\right) + \Theta(1)\right) + \Theta(1) \\ &= \left(\left(T\left(\frac{n}{8}\right) + \Theta(1)\right) + \Theta(1)\right) + \Theta(1) \end{aligned}$$

And so for k substitutions, we get

$$T(n) = T\left(\frac{n}{2^k}\right) + k\Theta(1)$$

And to get to our base case, we need $k = \log n$, getting

$$\begin{aligned} T(n) &= T(1) + (\log n)\Theta(1) \\ &= \Theta(1) + \log n \cdot \Theta(1) \\ &= \Theta(\log n) \end{aligned}$$

As expected. Always reassuring when new methods give us the answers we know are correct a priori!




Figure 1: A recursion tree for a call to MergeSort of size n

1.2 Recurrence Trees

Another method to analyze the runtime of recursive functions is to study a *recurrence tree*, a tree structure where each node represents a call to the function and its children denote the calls to its children. We annotate each node with the cost of executing that call *excluding work done by the recursive calls*. In that way, the cost of a call including the recursive call is the sum of the annotations of the corresponding node and all its descendants. Thus, the cost of calling the first, non-recursive call is the sum of the annotations to all nodes in the tree!

Take, for example, MERGESORT. We observe that in the first call, we do $\Theta(n)$ work outside of calls, and make two recursive calls, so we draw a tree with a root labeled $\Theta(n)$ to indicate the work done by the first call, and give it two children.

Those children will do $\Theta(\frac{n}{2})$ work (dominated by a call of size $\frac{n}{2}$ to MERGE), and each make two calls, and so on until a base case is reached. We can then identify a couple of features of the tree we can measure: the total amount of work done at each level of the tree (In this case, 2^k calls each doing $\Theta(\frac{n}{2^k})$ work), and the height of the tree (Since the tree splits the problem size in half each level, roughly $\log_2 n$ levels to get to our base case). This gives us Fig. 1.

From these features, we can compute the sum of the annotations on all nodes, summing level-by level

$$\begin{aligned} T(n) &= \sum_{k=0}^{\log n} 2^k \Theta\left(\frac{n}{2^k}\right) \\ &= \log n \cdot \Theta(n) = \Theta(n \log n) \end{aligned}$$

Converging with our prior result.

As you can see, both of these informal techniques converge on not only the same result, but the same eventual summation. The perk of seeing both is that they take different angles on what this the same mathematics: unrolling is a bit more directly math-y, and models the recursion via substitution. The recurrence trees help visualize the call structure and give you a way to see how the terms we gather in unrolling correspond to recursive depth in the call structure. The recurrence

trees also help us give language to dynamics that emerge in different classes of algorithms. These patterns are captured in what is called the *Master Theorem*

2 The Master Theorem

Theorem 2.1 (The Master Theorem) *Given a recurrence relation*

$$T(n) = \begin{cases} aT\left(\frac{n}{b}\right) + f(n), & n \geq k \\ \Theta(1), & n < k \end{cases}$$

where $a, b \in \mathbb{N}$,

1. If $f(n) \in \Theta(n^{\log_b a})$, then $T(n) \in \Theta(n^{\log_b a} \log n)$
2. If $f(n) \in O(n^c)$, with $c < \log_b a$, then $T(n) \in \Theta(n^{\log_b a})$
3. If $f(n) \in \Omega(n^c)$ with $c > \log_b a$ and $af\left(\frac{n}{b}\right) \leq kf(n)$ for $0 \leq k \leq 1$, then $T(n) = \Theta(f(n))$

To understand the Master Theorem, it's important to get a handle on the value $n^{\log_b a}$, which shows up in each case! This is the number of leaves (i.e., base cases!) in the recursive tree formed by a call of size n . Observe that this is actually just

$$\begin{aligned} n^{\log_b a} &= a^{\log_a (n^{\log_b a})} \\ &= a^{\log_b a \log_a n} \\ &= a^{\log_b n} \end{aligned}$$

Then note that a is the *branching factor* of the tree (how many children each node will have) and $\log_b n$ is the depth of our recursive tree!

Then we should note that since each base case takes $\Theta(1)$ time to complete, this means that just solving all of the base cases takes $\Theta(n^{\log_b a})$ time! The key dynamic here is whether our summation is dominated by the cost of solving the base cases, doing non-recursive work in each call, or whether these are balanced against each other.

There recurrence tree can help us give each of these cases a visual:

1. In Case 1, we have something that looks like MERGESORT! The amount of work done at each node ($f(n)$) is proportional to the amount of work done by all of the leaves in its subtree ($\Theta(n^{\log_b a})$), and so the amount of work done at each level will be proportional! In this case, our total runtime is the work done at each level ($\Theta(n^{\log_b a})$) multiplied by the number of levels ($\Theta(\log n)$)!
2. In Case 2, the work done at each node ($f(n)$) is *strictly*¹ less than the work done by the leaves in its subtree. This means that the work done in the tree is dominated by work done at the

¹Note that this is the purpose of c ! The exponent of n must be strictly smaller than $\log_b a$!

leaves! Thus, our runtime complexity is simply the cost of handling all of our base cases: $\Theta(n^{\log_b a})$!

3. Case 3 is a bit more complex. Here, we have that the work done at each node strictly dominates the work done by the leaves it generates. This alone gets us nowhere though, since this just tells us that the leaves won't dominate and nothing about how to sort out all of the internal layers! We need an additional assumption: That for any given node, the work done by its children ($af(\frac{n}{b})$) is only some fraction k of the work done at that node itself ($f(n)$). Together, this information tells us that the root's time complexity dominates: There aren't that many base cases, and the work at lower levels is just some fraction of work at a higher level, so our time complexity is just the work done at the highest level, $\Theta(f(n))$!